

## Problem Description

The problem I will be discussing involves parallelizing Merge Sort and counting the number of inversions. The original problem is provided in Polish in the file **instructions.md**.

## Tests

I have generated 5,000 test cases, each consisting of an array of size in the range from 20,000 to 100,000.

### Single-thread

I have introduced a few optimizations for the single-threaded case. The table below shows the average times for each program on several randomly generated tests:

|   | Avg Time (s) | Code Description                           |
|---|--------------|--|
| 1 | 20.39        | A naive merge sort implementation          |
| 2 | 9.72         | Disabled C/C++ stream synchronization      |
| 3 | 10.57        | Added insertion sort for small subproblems |
| 4 | 9.69         | Limited the number of comparisons          |

Unsurprisingly, adding *magic lines* speeds up the process immensely. However, employing insertion sort for small subproblems, which seemed like an obvious speed-up since it avoids the recursive stack, turned out to decrease performance. Finally, I limited the number of comparisons in the code, which resulted in a small speedup

Other potential optimizations could include using built-in functions (e.g., `memcpy` instead of manually copying arrays), avoiding repeated initialization of the `h` array by reusing a single buffer, and making better use of processor registers (e.g., incrementing a local counter in the while loop and updating the global counter only at the end).

### std::threads

Starting from a naive parallelization, I have introduced a few optimizations, including:

- introduced a local variable instead of updating the same counter with synchronization in the loop,
- introduced a threshold for parallelization: if a subproblem size is smaller than the threshold, no parallelization is used,
- introduced a shared buffer to avoid creating the same array repeatedly,
- counted inversions without using a global atomic counter.

|   | Avg Time (s) | Code Description                         |
|---|--------------|--|
| 1 | 11.24        | A naive parallelization                  |
| 2 | 7.66         | Optimized addressing shared resources    |
| 3 | 7.08         | Introduced a shared buffer               |
| 4 | 6.72         | Avoided using a global inversion counter |

## openMP

For a long time, I struggled to speed up OpenMP parallelization—on average, the results were equal to or worse than using a single thread. It turned out that calling `#pragma omp parallel` multiple times was the bottleneck, as it spawns multiple threads repeatedly. I managed to bypass this and ended up with the following results:

| x | Avg Time (s) | Code Description                      |
|---|--------------|---------------------------------------|
| 1 | 9.56         | A naive use of the library            |
| 2 | 6.17         | Optimized addressing shared resources |

## More threads

The tests above were run on my local machines. For each type, I selected one and ran tests on the student machine. Notice that beyond a certain point, execution does not get faster - as we optimized

|               | Threads | Avg Time (s) |
|---------------|---------|--------------|
| single thread | 1       | 12.04        |
| threads       | 8       | 9.31         |
| openmp        | 8       | 6.08         |
| threads       | 12      | 6.59         |
| openmp        | 12      | 6.15         |
| threads       | 16      | 7.25         |
| openmp        | 16      | 6.55         |
| threads       | 64      | 6.12         |
| openmp        | 54      | 6.46         |

only the recursion and not the merging step. Additionally, performance seems to get worse if too much threads are used. Presumably, this is due to the overhead of managing threads.

## Instructions

Running `python3 run.py` in the root folder generates five large tests (this may take a while, so you may want to change `TEST_CASES` from 1000 to 100), runs these C++ programs, measures the average execution time, and checks whether the results are correct.

In the `code` folder, you will find all the programs discussed above. You can modify the number of threads used. By default, it is set to 8.